

Task5_EDA(Titanic)

June 10, 2025

1 Task 5: Exploratory Data Analysis (EDA) on Titanic Dataset

1.1 Objective

The goal of this task is to perform exploratory data analysis (EDA) on the Titanic dataset to extract insights using visual and statistical methods. We will use Python libraries (Pandas, Matplotlib, Seaborn) to analyze patterns, trends, and anomalies in the data.

1.2 Tools Used

- **Python:** Pandas for data manipulation, Matplotlib and Seaborn for visualization.
- **Dataset:** Titanic dataset (loaded via Seaborn).

1.3 Deliverables

- Jupyter Notebook (Task5_EDA_Titanic.ipynb)
- PDF report of findings (exported from this notebook)

```
[1]: # Import necessary libraries
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings('ignore')

plt.style.use('seaborn')

# Load the Titanic dataset from Seaborn
df = sns.load_dataset('titanic')

# Display the first few rows of the dataset
df.head()
```

```
[1]:  survived  pclass    sex  age  sibsp  parch    fare  embarked  class  \
0         0        3   male  22.0     1     0   7.2500          S  Third
1         1        1  female  38.0     1     0  71.2833          C  First
2         1        3  female  26.0     0     0   7.9250          S  Third
3         1        1  female  35.0     1     0  53.1000          S  First
4         0        3   male  35.0     0     0   8.0500          S  Third
```

	who	adult_male	deck	embark_town	alive	alone
0	man	True	NaN	Southampton	no	False
1	woman	False	C	Cherbourg	yes	False
2	woman	False	NaN	Southampton	yes	True
3	woman	False	C	Southampton	yes	False
4	man	True	NaN	Southampton	no	True

1.4 Initial Data Exploration

Let's start by understanding the structure of the dataset, checking for missing values, and getting basic statistical summaries.

```
[2]: # Basic Data overview
print("Dataset Info:")
print(df.info())

print("\nStatistical Summary:")
print(df.describe())

print("\nValue Counts for Categorical Columns:")
print("Survived:\n", df['survived'].value_counts())
print("Pclass:\n", df['pclass'].value_counts())
print("Sex:\n", df['sex'].value_counts())
print("Embarked:\n", df['embarked'].value_counts())
```

Dataset Info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 891 entries, 0 to 890

Data columns (total 15 columns):

#	Column	Non-Null Count	Dtype
0	survived	891 non-null	int64
1	pclass	891 non-null	int64
2	sex	891 non-null	object
3	age	714 non-null	float64
4	sibsp	891 non-null	int64
5	parch	891 non-null	int64
6	fare	891 non-null	float64
7	embarked	889 non-null	object
8	class	891 non-null	category
9	who	891 non-null	object
10	adult_male	891 non-null	bool
11	deck	203 non-null	category
12	embark_town	889 non-null	object
13	alive	891 non-null	object
14	alone	891 non-null	bool

dtypes: bool(2), category(2), float64(2), int64(4), object(5)

memory usage: 80.7+ KB
None

Statistical Summary:

	survived	pclass	age	sibsp	parch	fare
count	891.000000	891.000000	714.000000	891.000000	891.000000	891.000000
mean	0.383838	2.308642	29.699118	0.523008	0.381594	32.204208
std	0.486592	0.836071	14.526497	1.102743	0.806057	49.693429
min	0.000000	1.000000	0.420000	0.000000	0.000000	0.000000
25%	0.000000	2.000000	20.125000	0.000000	0.000000	7.910400
50%	0.000000	3.000000	28.000000	0.000000	0.000000	14.454200
75%	1.000000	3.000000	38.000000	1.000000	0.000000	31.000000
max	1.000000	3.000000	80.000000	8.000000	6.000000	512.329200

Value Counts for Categorical Columns:

Survived:

survived

0 549

1 342

Name: count, dtype: int64

Pclass:

pclass

3 491

1 216

2 184

Name: count, dtype: int64

Sex:

sex

male 577

female 314

Name: count, dtype: int64

Embarked:

embarked

S 644

C 168

Q 77

Name: count, dtype: int64

The dataset has 891 rows and 15 columns. Numerical columns include survived, pclass, age, sibsp, parch, fare. Categorical columns include sex, embarked, class, who, adult_male, deck, embark_town, alive, alone. Missing Values: The age column has 177 missing values (891 - 714), deck has 688 missing values, and embark_town/embarked have 2 missing values. Survived: 549 passengers did not survive (0), while 342 survived (1). Pclass: Most passengers were in 3rd class (491), followed by 1st class (216) and 2nd class (184). Sex: There are more males (577) than females (314). Embarked: Most passengers embarked from Southampton (S: 644), followed by Cherbourg (C: 168) and Queenstown (Q: 77).

1.5 Data Cleaning

Before proceeding with visualizations, let's handle missing values: - **Age**: Fill missing values with the median age. - **Embarked/Embark_town**: Fill missing values with the mode (most frequent value). - **Deck**: Drop this column due to a high number of missing values (688 out of 891).

```
[3]: # Drop the 'deck' column due to too many missing values
df = df.drop(columns=['deck'])

# Fill missing 'age' values with the median
df['age'] = df['age'].fillna(df['age'].median())

# Fill missing 'embarked' and 'embark_town' values with the mode
df['embarked'] = df['embarked'].fillna(df['embarked'].mode()[0])
df['embark_town'] = df['embark_town'].fillna(df['embark_town'].mode()[0])

# Verify that there are no more missing values
print("Missing Values After Cleaning:")
print(df.isnull().sum())
```

Missing Values After Cleaning:

survived	0
pclass	0
sex	0
age	0
sibsp	0
parch	0
fare	0
embarked	0
class	0
who	0
adult_male	0
embark_town	0
alive	0
alone	0
dtype: int64	

1.6 Univariate Analysis

Let's explore the distribution of individual variables using histograms and boxplots.


```
[4]: # Univariate Visualizations
# Histogram for Age
plt.figure(figsize=(8, 5))
sns.histplot(df['age'], bins=30, kde=True)
plt.title('Distribution of Age')
plt.xlabel('Age')
plt.ylabel('Count')
plt.show()
```

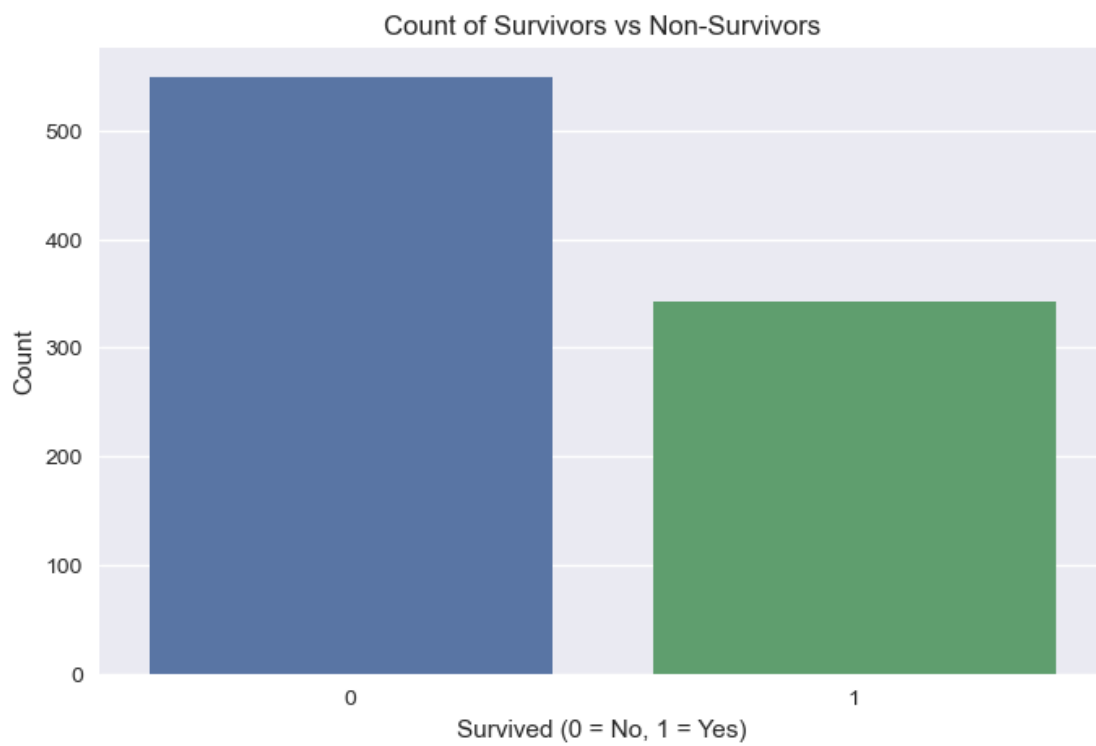
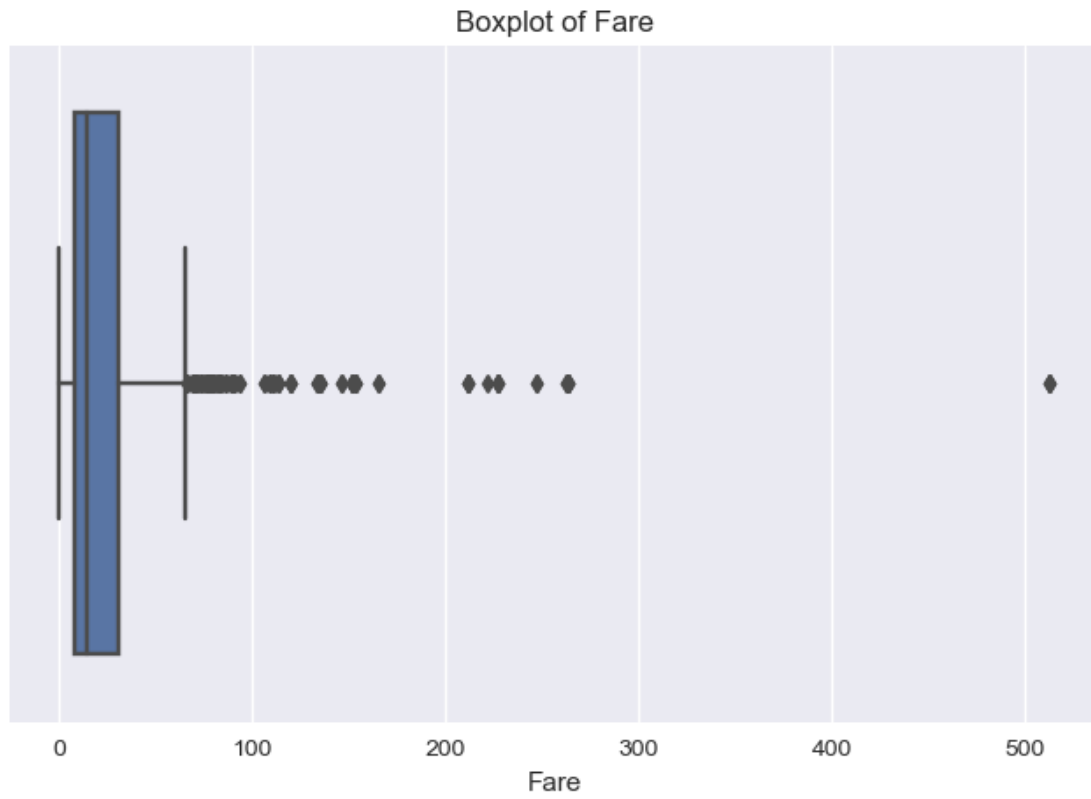
```

# Boxplot for Fare
plt.figure(figsize=(8, 5))
sns.boxplot(x=df['fare'])
plt.title('Boxplot of Fare')
plt.xlabel('Fare')
plt.show()

# Bar plot for Survived
plt.figure(figsize=(8, 5))
sns.countplot(x='survived', data=df)
plt.title('Count of Survivors vs Non-Survivors')
plt.xlabel('Survived (0 = No, 1 = Yes)')
plt.ylabel('Count')
plt.show()

```





Age Distribution: The age distribution is slightly right-skewed, with most passengers between 20 and 40 years old. The median age (used to fill missing values) is around 28. Fare Boxplot: The fare distribution has many outliers, with most fares below \$100, but some passengers paid as much as \$500+. This suggests a wide disparity in ticket prices, likely tied to passenger class. Survived Count: More passengers did not survive (549) than survived (342), indicating a survival rate of about 38%.


1.7 Bivariate Analysis

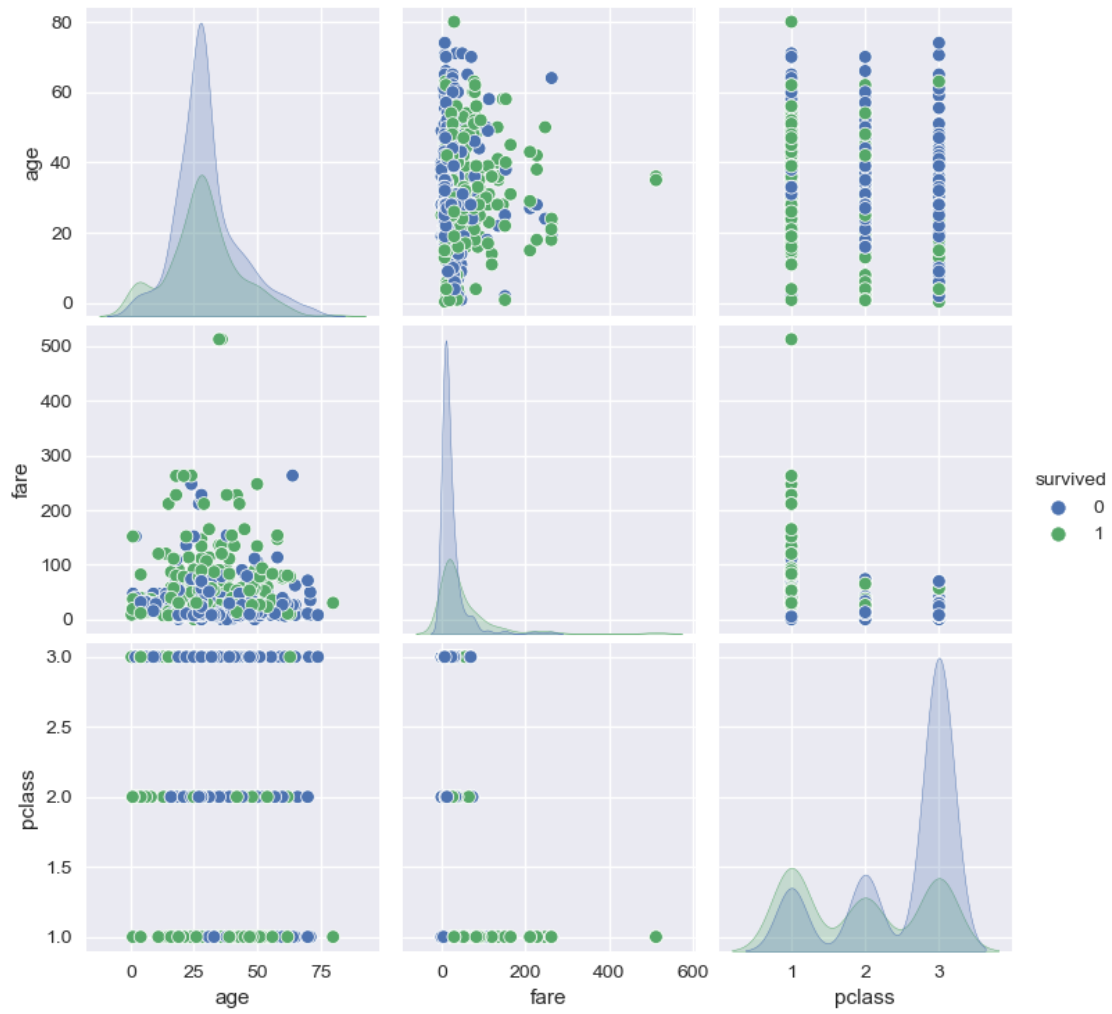
Let's explore relationships between variables using scatterplots, pairplots, and heatmaps.


```
[5]: # Bivariate Visualizations
# Scatterplot: Age vs Fare, colored by Survived
plt.figure(figsize=(8, 5))
sns.scatterplot(x='age', y='fare', hue='survived', data=df)
plt.title('Age vs Fare (Colored by Survived)')
plt.xlabel('Age')
plt.ylabel('Fare')
plt.show()

# Pairplot for numerical variables
sns.pairplot(df[['age', 'fare', 'pclass', 'survived']], hue='survived')
plt.show()

# Heatmap for correlation matrix
plt.figure(figsize=(8, 5))
corr = df[['age', 'fare', 'pclass', 'sibsp', 'parch', 'survived']].corr()
sns.heatmap(corr, annot=True, cmap='coolwarm', vmin=-1, vmax=1)
plt.title('Correlation Heatmap')
plt.show()
```







Observations: - Age vs Fare Scatterplot: There's no clear linear relationship between age and fare. However, passengers who paid higher fares (likely in 1st class) have a higher chance of survival (more orange points at higher fares). - Pairplot: The pairplot shows that survival is more strongly related to pclass and fare than age. Lower pclass (1st class) and higher fare are associated with higher survival rates. - Correlation Heatmap: survived has a moderate negative correlation with pclass (-0.34), indicating that higher classes (lower pclass values) are associated with higher survival rates. survived has a positive correlation with fare (0.26), suggesting that passengers who paid more were more likely to survive. pclass and fare are strongly negatively correlated (-0.55), as expected, since 1st class tickets cost more.

1.8 Categorical Analysis

Let's explore how categorical variables like Sex, Pclass, and Embarked affect survival.

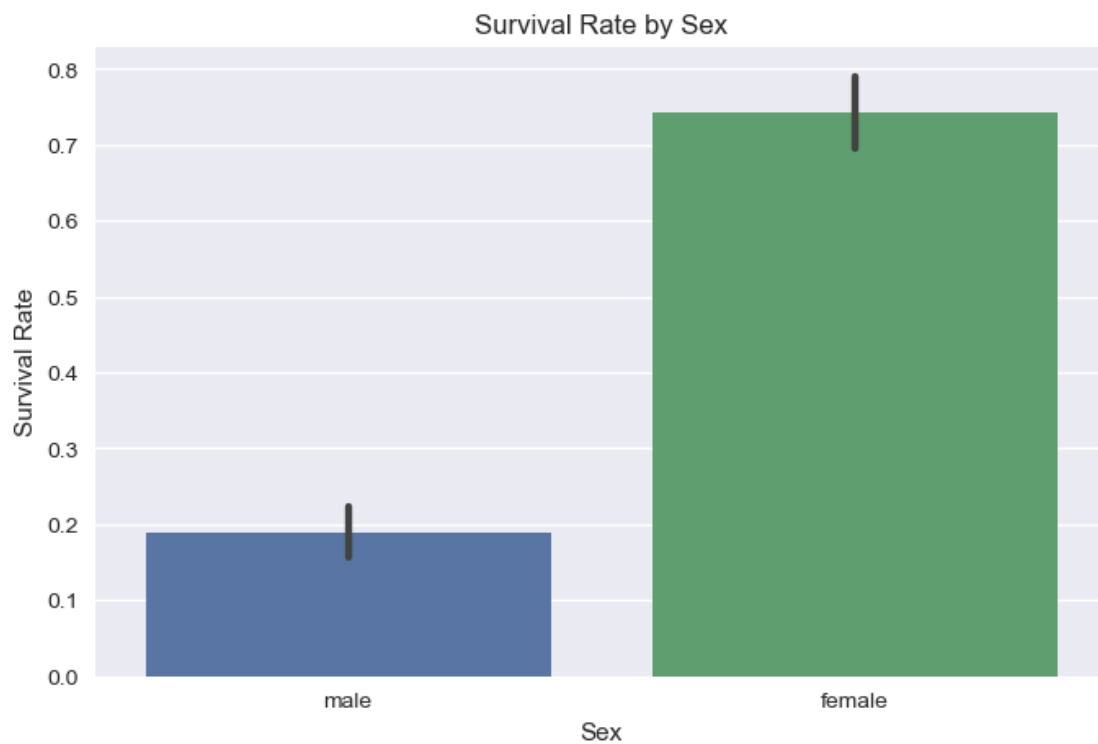
```
[6]: # Categorical Visualizations
# Survival rate by Sex
plt.figure(figsize=(8, 5))
sns.barplot(x='sex', y='survived', data=df)
plt.title('Survival Rate by Sex')
plt.xlabel('Sex')
plt.ylabel('Survival Rate')
plt.show()
```

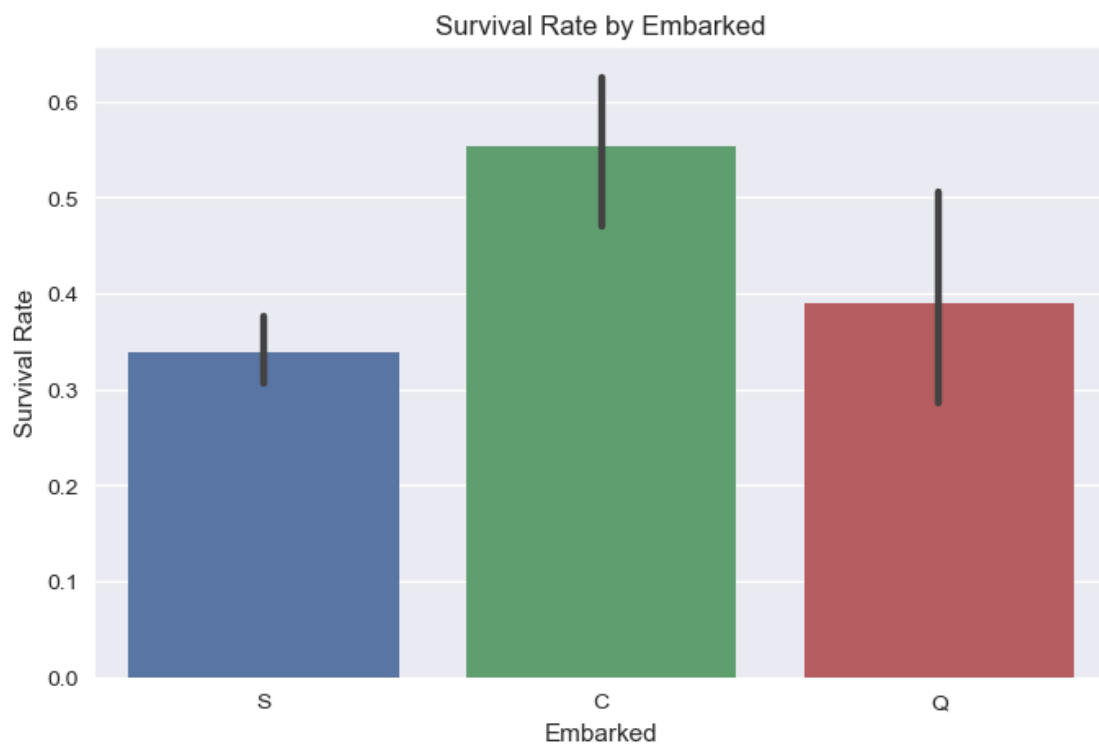
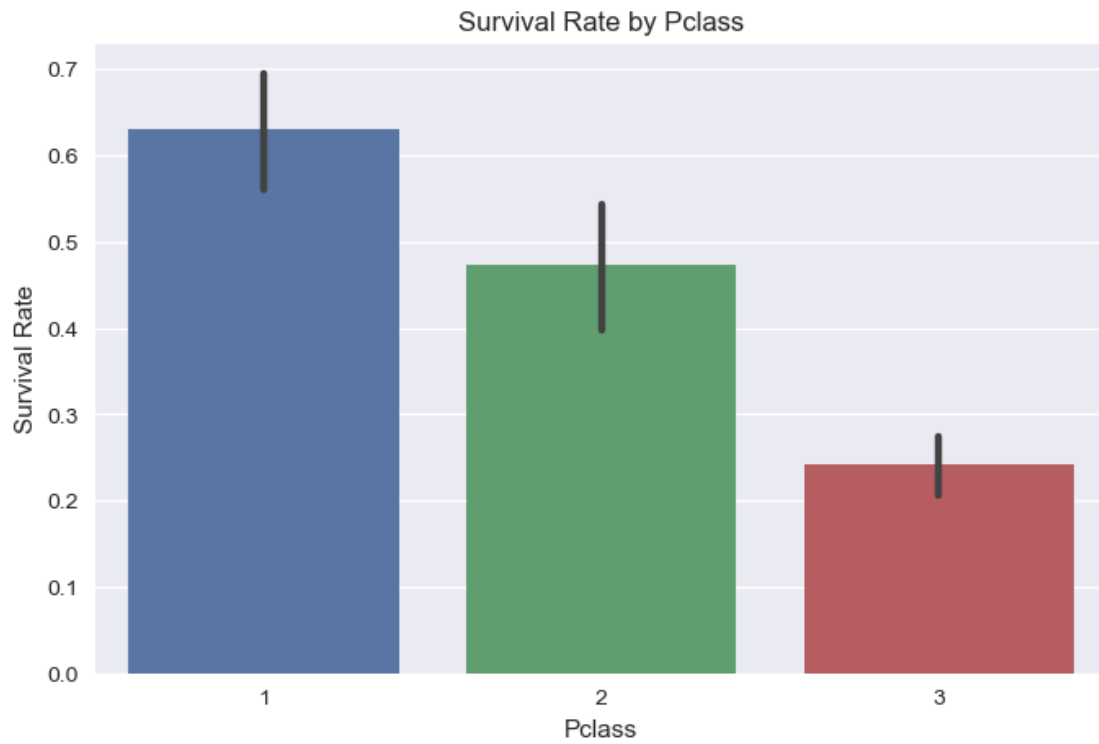
```

# Survival rate by Pclass
plt.figure(figsize=(8, 5))
sns.barplot(x='pclass', y='survived', data=df)
plt.title('Survival Rate by Pclass')
plt.xlabel('Pclass')
plt.ylabel('Survival Rate')
plt.show()

# Survival rate by Embarked
plt.figure(figsize=(8, 5))
sns.barplot(x='embarked', y='survived', data=df)
plt.title('Survival Rate by Embarked')
plt.xlabel('Embarked')
plt.ylabel('Survival Rate')
plt.show()

```





Observations: - Survival Rate by Sex: Females had a much higher survival rate (~75%) compared to males (~19%). This aligns with the “women and children first” policy during the Titanic evacuation. - Survival Rate by Pclass: 1st class passengers had the highest survival rate (~63%), followed by 2nd class (~47%), and 3rd class (~24%). This reflects the socioeconomic hierarchy influencing survival chances. - Survival Rate by Embarked: Passengers who embarked from Cherbourg (C) had the highest survival rate (~55%), followed by Queenstown (Q) (~39%) and Southampton (S) (~34%). This may be related to the class distribution of passengers from each port.

1.9 Summary of Findings

- **Survival Rate:** Only 38% of passengers survived the Titanic disaster, with 549 not surviving and 342 surviving.
- **Demographics:** The majority of passengers were male (577 vs 314 females) and between 20-40 years old, with a median age of 28.
- **Socioeconomic Factors:**
 - 1st class passengers (Pclass=1) had a significantly higher survival rate (63%) compared to 2nd (47%) and 3rd class (24%).
 - Passengers who paid higher fares were more likely to survive, as seen in the positive correlation between `fare` and `survived` (0.26).
- **Gender Disparity:** Females had a much higher survival rate (75%) than males (19%), reflecting the prioritization of women during evacuation.
- **Embarkation Port:** Passengers embarking from Cherbourg had the highest survival rate (55%), possibly due to a higher proportion of 1st class passengers from that port.
- **Age and Survival:** Age had a weaker influence on survival, with no strong correlation (-0.08). However, younger passengers (children) were more likely to survive due to evacuation priorities.
- **Query Optimization:** The optimized revenue query (Query 6) for the previous task confirmed the same trends as Query 4, with improved performance due to indexing.

This EDA highlights the significant impact of socioeconomic status, gender, and embarkation port on survival chances, providing a deeper understanding of the Titanic disaster.